

AI POWERED DRUG & DOSAGE RECOMMENDATION SYSTEM

A Project work

Submitted in partial fulfilment of the requirements for the award of the degree of

MASTER OF COMPUTER APPLICATIONS

By

VEMULA SYAM PRAKASH

Regd. No. 24FE1F00I3

Under the Esteemed Guidance of

Ms. K. Navya Blessi, MCA

Assistant Professor

DEPARTMENT OF MASTER OF COMPUTER APPLICATIONS



VIGNAN'S LARA
INSTITUTE OF TECHNOLOGY & SCIENCE
(AUTONOMOUS)

Approved by AICTE New Delhi & Affiliated to JNTUK Kakinada
Accredited by **NAAC 'A+' and NBA (CSE, IT, ECE, EEE & ME) | ISO 9001 : 2015**
Vadlamudi - 522 213, Guntur District

July - 2026



VIGNAN'S LARA

INSTITUTE OF TECHNOLOGY & SCIENCE

(AUTONOMOUS)

Approved by AICTE New Delhi & Affiliated to JNTUK Kakinada
Accredited by **NAAC 'A+' and NBA (CSE, IT, ECE, EEE & ME) | ISO 9001 : 2015**
Vadlamudi - 522 213, Guntur District

CERTIFICATE

This is to certify that the thesis entitled “**AI-POWERED DRUG AND DOSAGE RECOMMENDATION SYSTEM**” is a Bonafide work done by **Vemula Syam Prakash** with Regd No. 24FE1F00I3, in partial fulfilment for the award of Master of Computer Applications in Vignan’s Lara Institute of Technology & Science affiliated to Jawaharlal Nehru Technological University Kakinada, Kakinada is a record of Bonafide work carried out by him under my guidance and supervision. The work embedded in this project report is not submitted to any other university for the award of any degree/diploma.

Project Guide
Ms. Navya Blessi, MCA
Assistant Professor

Head of the Department
Mr. R. Veera Babu, M. Tech (Ph.D)
Associate Professor & HOD

External Examiner

DECLARATION

I hereby declare that the work described in this project work entitled “**AI-POWERED DRUG AND DOSAGE RECOMMENDATION SYSTEM**” submitted by me in partial fulfilment for the award of **Master of Computer Applications** in Vignan’s Lara Institute of Technology & Science, Vadlamudi, affiliated to Jawaharlal Nehru Technological University Kakinada, Kakinada, Andhra Pradesh, is the result of work done by me under the guidance of **Ms. K. Navya Blessi, MCA** Assistant Professor. The work is original and has not been submitted for any degree or diploma of this or any other university.

Place: Vadlamudi

Date:

V. Syam Prakash
(Regd. No. 24FE1F00I3)

ACKNOWLEDGMENT

The satisfaction that accompanies the successful completion of any task would be incomplete without the mention of the people whose ceaseless cooperation made it possible, and whose constant guidance and encouragement crowned all efforts with success.

I am glad to express my deep sense of gratitude to **Ms. K. NAVYA BLESSI**, Assistant professor, Master of Computer Applications for guiding through this project and for encouraging tight from the beginning of the project. Every interaction with her was an inspiration. At every step she was there to help me to choose right path

I am glad to express my deep sense of gratitude to **Mr. R. VEERA BABU** MTech (Ph.D.), Associate Professor and Head of the Department, for guiding me through this project and for encouraging me right from its beginning. Every interaction with him was an inspiration, and at every step he was there to help me choose the right path.

I also express my deep gratitude to our Principal, **Dr. K. PHANEENDRA KUMAR**, and to the beloved Chairman, **Dr. LAVU RATHAIAH**, for their encouragement and kind support in carrying out my work.

I thank my parents and classmates who rendered help to me, directly or indirectly, in the completion of this project.

VEMULA SYAM PRAKASH
(Regd. No. 24FE1F00I3)

ABSTRACT

The AI-Powered Drug and Dosage Recommendation System mark a significant advancement in the field of clinical decision support, addressing the growing complexity of drug selection and dosage optimization in modern healthcare. As the number of available pharmaceutical compounds increases exponentially, healthcare professionals face mounting challenges in selecting the most appropriate drug and dosage for individual patients, considering factors such as age, weight, medical history, symptoms, and drug interaction profiles. This project develops an intelligent recommendation system that leverages ensemble machine learning techniques, specifically Random Forest and Gradient Boosting classifiers, to provide accurate drug recommendations based on patient clinical data. The system integrates a comprehensive data preprocessing pipeline that handles missing values, categorical encoding, and feature scaling to ensure high-quality inputs to the machine learning models. Feature importance analysis using SHAP (Shapley Additive explanations) values is incorporated to enhance model interpretability and provide explainable AI-driven recommendations, bridging the gap between black-box predictions and clinically transparent decision support. The proposed system achieves a classification accuracy of over 92% on the UCI drug dataset, demonstrating its effectiveness in mapping patient clinical profiles to appropriate drug categories. The project is implemented as a web-based application using the Flask framework, deployed on Render.com, enabling healthcare professionals to input patient parameters through an intuitive interface and receive instant drug recommendations. Evaluation results show that the Random Forest classifier outperforms traditional rule-based systems, with a precision of 91%, recall of 89%, and an F1-score of 90%. The system's ability to process patient data in real time and deliver explainable recommendations makes it a valuable tool for assisting clinicians in personalized medicine.

INDEX

CONTENTS	PAGE NO
CERTIFICATE	I
DECLARATION	II
ACKNOWLEDGEMENT	III
ABSTRACT	IV
Chapter 1 INTRODUCTION	1-9
1.1 Introduction	1
1.2 Purpose	2
1.3 Scope	2
1.3.1 Architectural Scope and Data Boundaries	3
1.3.2 Ethical and Regulatory Scope Boundary	3
1.4 Objective	3-4
1.5 Importance of Artificial Intelligence in Healthcare	5
1.6 Challenges in Drug Recommendation	6-8
1.7 Benefits of AI-Powered Drug Recommendation System	9
Chapter 2 LITERATURE REVIEW	10-13
2.1 Literature Survey	10-13
Chapter 3 SYSTEM ANALYSIS	14-27
3.1 Existing System	14-15
3.1.1 Limitations of Existing Systems	16
3.2 Proposed System	17
3.2.1 Feasibility Study	18
3.2.2 Advanced Clinical “What-If” Simulation Module	19
3.2.3 Predictive Side-Effect Analysis	19
3.2.4 Consensus Tiebreaking Mechanism	20
3.2.5 Dosage Tier Determination Logic	20
3.3 Data Collection	21
3.3.1 Workflow	21-22
3.3.2 Data Processing	23
3.4 Minimum Hardware Requirements	24
3.5 Minimum Software Requirements	25

3.5.1 Functional Requirements	25
3.5.2 Non-Functional Requirements	25-27
Chapter 4 SYSTEM DESIGN	28-40
4.1 System Design	28
4.1.1 High Level System Architecture	28
4.1.2 Key Components & Their Interactions	28
4.1.3 Data Flow & Communication Pathways	29
4.1.4 UML Diagrams	29-30
4.1.5 Mathematical Foundations of the Inference Engine	31
4.2 Implementation	32
4.2.1 Technical Stack	32
4.2.2 Code Samples	33-34
4.3 Testing	35
4.3.1 Testing Introduction	35
4.3.2 Testing Process	35
4.3.3 Levels of Testing	36-40
Chapter 5 RESULT & DISCUSSION	41-51
5.1 Result	41
5.1.1 Per-Class Performance Analysis	42
5.1.2 Cross-Validation Stability Analysis	43
5.2 Discussion	43-49
5.3 Simulation and Side-Effect Evaluation	50-51
Chapter 6 CONCLUSION & FUTURE WORK	52-57
6.1 Research Limitations and Mitigation Pathways	55
6.2 Clinical Deployment Roadmap	55-57
REFERENCES	58

LIST OF FIGURES

Figure No	Figure Name	Page No
Fig 1.7.1	Benefits of AI-Powered Drug Recommendation System	9
Fig 3.1.1	Traditional Drug Prescription Workflow	15
Fig 3.2.1	Proposed AI Drug Recommendation System Architecture	17
Fig 3.3.2.1	Data Processing Pipeline	24
Fig 4.1.1.1	High-Level System Architecture	28
Fig 4.1.4.1.1	Use Case Diagram	30
Fig 4.1.4.2.1	Class Diagram	30
Fig 4.1.4.3.1	Sequence Diagram	30
Fig 4.3.2.1	Testing Process	35
Fig 4.3.3.1	Levels of Testing	37
Fig 5.2.1.1	Accuracy & Precision Score for Dataset	44
Fig 5.2.2.1	Recall and F1-Score Charts	45
Fig 5.2.3.1	Confusion Matrix	45
Fig 5.2.4.1	Bar Graphs for Comparative Analysis	46

LIST OF TABLES

Table	Table name	Page No
Table 3.3.1.1	Hyperparameter Search Grid for Ensemble Classifiers	22
Table 3.3.2.1	Descriptive Statistics of Continuous Clinical Features	23
Table 3.4.1	Minimum Hardware Requirements	24
Table 3.5.1	Minimum Software Requirements	25
Table 4.3.3.3.1	Performance Metrics	39
Table 5.2.6.1	McNemar's Test Contingency Table	48
Table 5.2.7.1	Global SHAP Feature Importance Ranking	48

CHAPTER 1

INTRODUCTION

1.1 INTRODUCTION

Healthcare has become increasingly dependent on accurate and timely clinical decisions to ensure effective patient treatment. One of the most important aspects of clinical practice is selecting the appropriate medication and dosage for each patient. This process is often influenced by several factors, including age, gender, blood pressure, cholesterol level, medical history, existing diseases, allergies, and laboratory findings. As the number of available medicines continues to grow, identifying the most suitable treatment has become more complex for healthcare professionals.

Medication-related errors remain a major concern worldwide because inappropriate drug selection or incorrect dosage can lead to ineffective treatment, adverse drug reactions, longer hospital stays, and increased healthcare costs. Conventional prescription methods mainly rely on physician knowledge, clinical experience, and standard treatment guidelines. Although these methods are effective, they may not always consider the unique characteristics of every patient or rapidly analyze large amounts of clinical data. Recent developments in Artificial Intelligence (AI) and Machine Learning (ML) have introduced new possibilities for improving clinical decision-making. Machine learning algorithms can analyse historical patient data, recognize hidden relationships among clinical variables, and generate reliable predictions that support medical professionals during treatment planning. Instead of replacing physicians, these intelligent systems function as decision-support tools that assist in selecting appropriate medications based on patient-specific information.

The proposed AI-Powered Drug and Dosage Recommendation System utilize ensemble machine learning algorithms, namely Random Forest and Gradient Boosting, to recommend suitable drugs using patient clinical parameters. The system performs comprehensive data preprocessing, feature transformation, and prediction generation before providing recommendations through a web-based interface. To improve transparency and clinician confidence, the model also incorporates SHAP (SHapley Additive Explanations), which explains how each patient feature contributes to the final prediction. This combination of prediction accuracy and explainability makes the proposed system a practical solution for supporting personalized healthcare.

1.2 PURPOSE

The healthcare industry continues to grapple with the challenge of medication errors, which arise from incorrect drug selection, inappropriate dosage, or failure to account for patient-specific contraindications. Traditional prescription practices often rely heavily on physician expertise and experience, which, while invaluable, can be inconsistent across practitioners and susceptible to cognitive overload in high-pressure clinical environments. The sheer volume of available pharmacological agents, each with unique mechanisms of action, side-effect profiles, and interaction risks, makes it increasingly difficult for clinicians to make optimal decisions without computational assistance.

Furthermore, the emergence of personalized medicine has underscored the need for recommendation systems that can account for individual patient variability. Factors such as Na/K (sodium-to-potassium) ratio, blood pressure, cholesterol levels, age, and sex significantly influence drug efficacy and safety profiles. Conventional drug selection protocols often apply generalized guidelines that may not adequately reflect the nuanced requirements of specific patient subgroups. This limitation highlights the urgent need for intelligent systems capable of processing multidimensional patient data and generating tailored recommendations.

The purpose of this project is to develop a machine learning-based drug recommendation system that bridges the gap between raw clinical data and actionable therapeutic guidance. By training predictive models on annotated patient datasets, the system learns to associate specific combinations of clinical features with optimal drug choices, enabling rapid and accurate recommendations that complement the clinical judgment of healthcare professionals. As highlighted by Chen et al. (2020), the integration of machine learning into clinical workflows provides a promising approach for strengthening drug prescription mechanisms and reducing medication errors.

1.3 SCOPE

The scope of this project is limited to developing a machine learning-based web application capable of recommending appropriate drug categories using structured patient information. The system accepts patient parameters such as age, sex, blood pressure, cholesterol level, and Na/K ratio as input and predicts the most suitable drug class based on patterns learned from the UCI Drug dataset. The application is developed using Python,

Flask, Scikit-learn, and SHAP, while deployment is carried out using the Render cloud platform. The system emphasizes prediction accuracy, model interpretability, and user-friendly interaction, enabling healthcare professionals to obtain recommendations through a simple web interface. However, the current implementation has certain limitations. It does not diagnose diseases, prescribe exact dosage quantities for individual patients, replace physician judgment, or integrate with Electronic Health Record (EHR) systems. These features can be incorporated in future versions of the project.

1.3.1 Architectural Scope and Data Boundaries

The operational boundary of the system is strictly defined by the clinical features available within the UCI machine learning dataset. The system acts exclusively as a downstream inference framework, meaning it consumes pre-extracted biochemical data (Na/K ratio) and vital indicators (Blood Pressure and Cholesterol) rather than directly interfacing with raw laboratory telemetry. Furthermore, the system assumes a static data environment per request; real-time continuous telemetry monitoring or historical trend cross-referencing over multiple longitudinal patient visits falls outside the immediate scope of this deployment stage.

1.3.2 Ethical and Regulatory Scope Boundary

Beyond its technical boundaries, the system is deliberately scoped to function as a clinician-facing decision-support aid rather than an autonomous prescribing agent. Every recommendation generated by the platform is presented as an advisory output accompanied by its SHAP-based explanation, and the interface makes no attempt to finalize, transmit, or dispense a prescription on its own. This design choice reflects a broader principle in clinical AI development: systems that influence treatment decisions must keep a qualified human reviewer in the decision loop, both to catch edge cases that fall outside the training distribution and to retain clear accountability for the final clinical action. Consequently, the current version of the system is positioned for academic demonstration and research evaluation rather than for unsupervised use in live patient care, and any future transition toward clinical deployment would first require the regulatory validation steps outlined later in the Clinical Deployment Roadmap.

1.4 OBJECTIVE

The structure of this report is meticulously organized to provide a comprehensive overview of the AI-Powered Drug and Dosage Recommendation System, detailing its development, implementation, and potential impact on the field of clinical decision support. The report begins with a detailed literature survey examining existing research on machine learning applications in clinical decision support and drug recommendation systems, as well as an analysis of current drug prescription frameworks and their limitations. This section sets the stage for understanding the gaps in current approaches that the proposed system aims to address. Following the literature survey, the system analysis section is divided into two subsections: the existing system and the proposed system. The existing system subsection outlines the methodologies employed by current drug prescription frameworks, highlighting their limitations and challenges. In contrast, the proposed system subsection introduces the unique features of the AI Drug Recommendation System, demonstrating how it addresses these limitations through the integration of ensemble machine learning and explainable AI techniques. The feasibility study section evaluates the technical, economic, and operational feasibility of implementing the system, assessing potential risks and proposing mitigation strategies. This is followed by the system requirement specifications, which detail the hardware and software requirements necessary for deploying the system, including scalability considerations.

The methodology section provides an in-depth examination of the workflow involved in drug classification, emphasizing key stages such as data collection, preprocessing, model training, and prediction generation. The testing section outlines the testing methodologies applied, including unit, integration, and system testing phases. The results and discussion section presents evaluation metrics results. A secondary objective of this project is to contribute a replicable methodological template for building explainable clinical AI systems in resource-constrained academic and institutional environments. By demonstrating that a high-accuracy, clinically interpretable drug recommendation system can be built using entirely open-source tools — scikit-learn, Flask, SHAP, and Render.com’s free tier cloud hosting — this project establishes a low barrier pathway for clinical informatics researchers and health technology developers in lower-resource settings to build and deploy similar systems. This democratization of clinical AI tooling represents an important step toward closing the global health technology equity gap, where innovative clinical decision support tools are currently concentrated in well-funded

healthcare systems in a small number of high-income countries. A tertiary objective is to generate empirical evidence that informs the ongoing academic debate around the relative clinical value of model accuracy versus model explainability in healthcare AI. The XAI literature includes perspectives that frame accuracy and explainability as inherently conflicting objectives — arguing that more accurate models are necessarily more complex and thus less interpretable. This project challenges that framing by demonstrating that tree-based ensemble models, which achieve competitive accuracy on clinical tabular datasets, are simultaneously amenable to exact SHAP value computation through the Tree Explainer algorithm, producing both high predictive performance and feature-level transparency without architectural compromise. The empirical results presented in this report therefore contribute evidence to the position that accuracy explainability trade-offs, while real in some model families, are not an inevitable constraint in well-scoped clinical classification tasks using structured patient data.

1.5 IMPORTANCE OF ARTIFICIAL INTELLIGENCE IN HEALTHCARE

Artificial Intelligence (AI) has emerged as one of the most transformative technologies in modern healthcare, enabling healthcare professionals to improve the quality, efficiency, and accessibility of medical services. AI systems utilize advanced computational techniques such as Machine Learning (ML), Deep Learning (DL), Natural Language Processing (NLP), and Computer Vision to analyze large volumes of healthcare data and generate intelligent recommendations. These technologies assist medical practitioners in making accurate clinical decisions while reducing the possibility of human error.

The increasing availability of electronic health records, laboratory reports, medical imaging, genomic information, and wearable health monitoring devices has resulted in an enormous volume of healthcare data. Traditional analytical methods are often unable to process such vast datasets efficiently. Artificial Intelligence provides the capability to analyse structured and unstructured medical data rapidly, discovering hidden patterns and relationships that support early diagnosis, disease prediction, and treatment planning.

One of the most significant applications of AI is Clinical Decision Support Systems (CDSS). These systems assist physicians by providing evidence-based recommendations during diagnosis, treatment selection, and medication prescribing. Instead of replacing healthcare professionals, AI serves as an intelligent assistant that enhances clinical

judgment by presenting relevant information based on patient-specific characteristics. Artificial Intelligence has also revolutionized personalized medicine. Conventional treatment methods generally follow standardized clinical guidelines that may not consider the unique biological characteristics of each patient. AI enables personalized treatment recommendations by analysing demographic information, clinical symptoms, laboratory parameters, medical history, and genetic information. As a result, treatment plans become more accurate, effective, and safer for individual patients.

Drug recommendation represents another rapidly growing application of AI in healthcare. Selecting the appropriate medication requires consideration of numerous factors, including patient age, gender, blood pressure, cholesterol levels, allergies, existing medical conditions, and possible drug interactions. AI models can simultaneously evaluate these variables and recommend suitable medications based on previously learned clinical patterns.

Machine learning algorithms have demonstrated remarkable performance in various medical domains, including disease diagnosis, drug discovery, cancer detection, medical image interpretation, patient monitoring, and healthcare resource management. Ensemble learning algorithms such as Random Forest and Gradient Boosting have gained particular importance because of their ability to achieve high prediction accuracy while reducing overfitting. These characteristics make them especially suitable for healthcare applications where reliability and consistency are essential.

The adoption of Explainable Artificial Intelligence (XAI) has further strengthened the practical applicability of AI in healthcare. Medical professionals require transparent reasoning behind AI-generated recommendations before incorporating them into clinical practice. Explainability techniques such as SHAP (SHapley Additive Explanations) provide detailed insights into how each patient feature contributes to the final prediction, thereby increasing clinician confidence and supporting informed decision-making.

Overall, Artificial Intelligence is transforming healthcare by enabling faster diagnosis, reducing medication errors, supporting personalized treatment, improving patient outcomes, and assisting healthcare professionals in making evidence-based clinical decisions. As AI technologies continue to evolve, they are expected to become indispensable components of future healthcare systems.

1.6 CHALLENGES IN DRUG RECOMMENDATION

Drug recommendation is one of the most complex tasks in clinical practice because it requires careful evaluation of multiple patient-specific factors. Healthcare professionals must consider medical history, physiological conditions, laboratory findings, allergies, concurrent medications, chronic diseases, and potential adverse drug reactions before selecting an appropriate treatment. Even experienced physicians may encounter challenges when dealing with complex cases involving multiple interacting variables. Medication errors remain a major global healthcare concern. According to international healthcare reports, medication errors contribute significantly to preventable patient injuries and healthcare costs. These errors may occur during prescription, dosage calculation, dispensing, or administration of medications. Incorrect drug selection or inappropriate dosage can result in severe adverse reactions, prolonged hospitalization, or even life-threatening complications.

Another major challenge is polypharmacy, which refers to the simultaneous use of multiple medications by a single patient. Polypharmacy is particularly common among elderly patients and individuals suffering from chronic diseases such as diabetes, hypertension, and cardiovascular disorders. The increasing number of prescribed medications significantly raises the risk of drug-drug interactions, adverse effects, and therapeutic conflicts. Patient diversity also complicates drug recommendation. Individuals with similar symptoms may require different medications depending on factors such as age, sex, body weight, kidney function, liver function, blood pressure, cholesterol levels, and metabolic characteristics. Standard treatment guidelines cannot always accommodate these individual differences, highlighting the need for intelligent decision support systems capable of personalized recommendations. Clinical decision-making often occurs under significant time constraints, particularly in emergency departments and high-volume healthcare facilities. Physicians must analyse large amounts of patient information within limited timeframes while maintaining high standards of accuracy. This cognitive burden increases the possibility of overlooking important clinical information, emphasizing the importance of AI-assisted recommendation systems. Data quality presents another significant challenge. Healthcare datasets frequently contain missing values, inconsistent measurements, duplicate records, and heterogeneous data formats collected from different clinical institutions. Effective preprocessing techniques, including data cleaning, feature

encoding, normalization, and validation, are essential before applying machine learning algorithms.

Model interpretability represents an additional concern. Many advanced machine learning models function as “black boxes,” providing highly accurate predictions without explaining the reasoning behind their decisions. In healthcare environments, clinicians require transparent and interpretable recommendations to ensure patient safety and regulatory compliance. Explainable AI techniques have therefore become increasingly important for bridging the gap between predictive accuracy and clinical trust. Finally, ethical and legal considerations must be addressed during AI implementation. Patient privacy, data security, fairness, algorithmic bias, and regulatory compliance remain essential aspects of healthcare AI deployment. Intelligent systems should support healthcare professionals rather than replace clinical expertise, ensuring that final treatment decisions remain under physician supervision.

Emerging Challenges in AI-Based Drug Recommendation

Another significant challenge is the continuous evolution of medical knowledge. New drugs, updated treatment guidelines, and newly discovered drug interactions are introduced regularly. Machine learning models trained on historical datasets may become outdated if they are not periodically retrained with recent clinical data. Therefore, continuous model updating and validation are essential to ensure that AI systems remain accurate and clinically relevant.

Another concern is the limited availability of high-quality healthcare datasets. Medical data often originate from different hospitals and healthcare organizations that use different data formats, coding standards, and documentation practices. These inconsistencies make data integration difficult and may reduce the overall performance of predictive models. Developing standardized healthcare data formats is therefore important for improving AI-based recommendation systems.

The adoption of Artificial Intelligence in healthcare also depends heavily on user acceptance. Many healthcare professionals remain cautious about relying on automated recommendations, especially when the underlying decision-making process is not fully transparent. Explainable Artificial Intelligence (XAI) techniques, such as SHAP, help address this issue by providing understandable explanations for each prediction, thereby

increasing clinician confidence and encouraging responsible adoption. Finally, ensuring continuous monitoring and evaluation of deployed AI systems is critical. Model performance may gradually decline due to changes in patient populations, disease patterns, or clinical practices, a phenomenon known as model drift. Regular performance evaluation, retraining, and validation are therefore necessary to maintain prediction accuracy, improve patient safety, and ensure compliance with evolving healthcare regulations. Scalability is another practical challenge.

1.7 BENEFITS OF AI-POWERED DRUG RECOMMENDATION SYSTEM




Fig 1.7.1: Benefits of AI-Powered Drug Recommendation System

As healthcare organizations generate millions of patient records, recommendation systems must efficiently process large volumes of data while maintaining low response times. Cloud computing technologies, distributed processing frameworks, and optimized machine learning algorithms are essential for supporting large-scale deployment in hospitals and healthcare networks. The system employs ensemble machine learning to analyse patient clinical data and generate intelligent drug recommendations that assist healthcare professionals in making accurate and timely treatment decisions. Collectively, these benefits lead to the final outcome of Improved Patient Outcomes. By combining high prediction accuracy, rapid processing, personalized recommendations, explainable artificial intelligence, and enhanced patient safety, the proposed AI-powered drug recommendation system serves as an effective clinical decision support tool that improves the overall quality, efficiency, and reliability of healthcare services.

CHAPTER 2

LITERATURE REVIEW

2.1 LITERATURE SURVEY

Artificial Intelligence (AI) and Machine Learning (ML) have significantly transformed healthcare by supporting clinicians in diagnosis, treatment planning, disease prediction, and medication recommendation. With the increasing availability of digital healthcare records and computational resources, researchers have developed intelligent systems capable of analysing patient information and generating accurate clinical recommendations. This chapter reviews major research contributions related to AI-driven drug recommendation systems, ensemble learning algorithms, and Explainable Artificial Intelligence (XAI), which form the foundation of the proposed system.

One of the pioneering studies in AI-assisted healthcare was conducted by Rajkumar et al. (2018), who demonstrated that deep learning models trained on large-scale Electronic Health Records (EHRs) could achieve high prediction accuracy for several clinical tasks, including mortality prediction, hospital readmission, and length of hospital stay. Their research established that machine learning models can effectively extract meaningful clinical patterns from complex healthcare datasets and provide valuable support to medical professionals.

Chen et al. (2020) investigated the application of machine learning techniques for drug recommendation using patient clinical characteristics. Their study compared several classification algorithms and reported that ensemble methods, particularly Random Forest and Gradient Boosting, produced more reliable predictions than conventional single-model classifiers. Their findings highlighted the importance of appropriate feature engineering and demonstrated that clinical parameters such as sodium-to-potassium ratio, blood pressure, and cholesterol significantly influence drug selection.

Another important contribution was presented by Garg et al. (2021), who proposed an intelligent framework for predicting suitable drug dosage levels using ensemble learning techniques. Their research indicated that combining multiple machine learning algorithms improved prediction accuracy while reducing model bias and variance. The study also emphasized that patient physiological characteristics should play a central role in determining appropriate medication dosage.

Explainability has become an essential requirement for medical AI systems. Zhang et al. (2022) introduced a SHAP-based explanation framework for clinical drug recommendation. Their work demonstrated that feature contribution analysis allows healthcare professionals to understand why a particular recommendation is generated by the model. Such transparent decision-making increases clinician confidence and supports the safe adoption of AI in healthcare environments.

A comparative evaluation conducted by Kumar et al. (2021) examined several supervised learning algorithms, including Naïve Bayes, Decision Tree, Support Vector Machine, Random Forest, and Gradient Boosting. Their experiments showed that ensemble learning methods consistently outperformed individual classifiers across multiple evaluation metrics such as accuracy, precision, recall, and F1-score. The study concluded that ensemble models provide better generalization capability and are less susceptible to overfitting, making them suitable for healthcare applications.

The theoretical basis for explainable machine learning was established by Lundberg and Lee (2017) through the introduction of SHAP (SHapley Additive Explanations). Their work applies concepts from cooperative game theory to determine the contribution of each feature toward a prediction. SHAP has become one of the most widely accepted explainability techniques because it generates consistent and mathematically reliable feature importance values, enabling clinicians to interpret AI-generated recommendations effectively.

The Random Forest algorithm, introduced by Breiman (2001), remains one of the most successful ensemble learning methods. By combining predictions from multiple decision trees created using random sampling techniques, Random Forest improves prediction stability and minimizes overfitting. Its ability to process nonlinear relationships and handle high-dimensional data makes it highly suitable for medical classification problems.

Similarly, Friedman (2001) proposed the Gradient Boosting Machine, which constructs predictive models by sequentially minimizing previous prediction errors. Unlike Random Forest, Gradient Boosting focuses on reducing bias through iterative optimization. The complementary strengths of these two algorithms make their combination particularly effective for healthcare prediction tasks requiring high reliability.

Ethical considerations in healthcare AI were highlighted by Obermeyer et al. (2019), who demonstrated that biased training data can produce unfair clinical recommendations. Their findings emphasize the importance of fairness evaluation, continuous monitoring, and transparent model development when deploying AI systems in healthcare environments.

In addition, Topol (2019) argued that Artificial Intelligence should function as an intelligent assistant rather than replacing healthcare professionals. According to his research, AI systems are most effective when they enhance physician decision-making by providing evidence-based recommendations while leaving the final clinical judgment to qualified medical practitioners.

Overall, previous research confirms that machine learning has considerable potential for improving drug recommendation and clinical decision support. However, many existing systems either lack sufficient prediction accuracy or fail to provide understandable explanations for their recommendations. The integration of ensemble learning with Explainable Artificial Intelligence addresses these limitations by delivering both accurate predictions and transparent decision-making.

Research Gap Synthesis

Read collectively, the studies reviewed above point to three recurring gaps that the present work is positioned to address. First, accuracy-focused studies such as Chen et al. (2020), Garg et al. (2021), and Kumar et al. (2021) consistently show that ensemble methods outperform single-model classifiers, yet none of these works pair their ensemble architecture with a rigorous, model-native explainability mechanism — explanations, where offered, are typically post-hoc and approximate rather than exact. Second, the explainability-focused literature, represented by Lundberg and Lee (2017) and Zhang et al. (2022), establishes the theoretical and practical value of SHAP but does not evaluate it in combination with a dual-ensemble consensus architecture of the kind proposed here, leaving open the question of whether combining two structurally distinct ensembles complicates or preserves the clarity of feature-level explanations. Third, the fairness and deployment literature, exemplified by Obermeyer et al. (2019) and Topol (2019), raises legitimate concerns about bias and clinician trust but stops short of demonstrating a concrete, reproducible pipeline that a resource-constrained academic or clinical team could adopt directly. The present project addresses all three gaps simultaneously: it combines

Random Forest and Gradient Boosting in a consensus framework (addressing the first gap), applies exact tree Explainer-based SHAP values to that consensus framework rather than to a single model in isolation (addressing the second gap), and documents the complete open-source implementation and deployment pipeline in sufficient detail to be replicated by other researchers (addressing the third gap).

Regulatory and Adoption Context

Beyond the algorithmic literature, the broader regulatory environment for AI-based clinical software has continued to evolve and forms an important part of the context in which systems of this kind must eventually operate. Regulatory bodies such as the United States Food and Drug Administration have progressively developed guidance for AI- and machine-learning-based Software as a Medical Device, emphasizing transparency, predetermined change control, and post-market performance monitoring as prerequisites for clinical deployment. International health authorities, including the World Health Organization, have similarly issued guidance stressing human oversight, fairness, and explainability as ethical prerequisites for deploying AI in health systems. These developments reinforce the design philosophy adopted in this project — namely, that a clinically useful drug recommendation system must treat explainability not as an optional add-on but as a core architectural requirement from the outset. This regulatory context also explains why the feasibility and deployment-roadmap discussions later in this report treat SHAP-based transparency and human-in-the-loop review as prerequisites for any future clinical pilot, rather than as features that could be retrofitted after the fact.

The proposed AI-Powered Drug and Dosage Recommendation System build upon these research contributions by combining Random Forest, Gradient Boosting, and SHAP explainability within a web-based clinical decision support platform. This approach aims to improve prediction performance, increase clinician trust, and provide a practical solution for personalized medication recommendation.

CHAPTER 3

SYSTEM ANALYSIS

3.1 EXISTING SYSTEM

Healthcare organizations have traditionally relied on manual prescription methods and conventional Clinical Decision Support Systems (CDSS) to assist physicians in selecting appropriate medications. These systems primarily use predefined clinical rules, treatment guidelines, and expert knowledge to recommend drugs based on patient conditions. Although such approaches have improved healthcare delivery over the years, they are often limited in their ability to process complex patient information and adapt to changing medical knowledge. Rule-based systems generate recommendations using fixed logical conditions. They compare patient information such as age, allergies, blood pressure, and laboratory findings against stored medical rules before suggesting suitable medications. Since these systems depend on manually created rules, they require frequent updates whenever new drugs or treatment guidelines become available.

Rule-Based Clinical Decision Support Systems (CDSS)

The earliest and most widely deployed drug recommendation systems in clinical practice are rule-based CDSS, which encode pharmacological knowledge as explicit IF-THEN decision trees and lookup tables derived from clinical guidelines. Systems such as First Databank, Micromedex, and Lexi-Comp operate by matching patient parameters typically diagnoses, allergies, and renal function indicators against curated drug formularies to generate recommendation alerts and contraindication warnings.




Fig 3.1.1: Traditional Drug Prescription Workflow

These systems offer significant advantages in regulatory compliance and clinical accountability: their decision logic is fully transparent, auditable, and directly traceable to published clinical guidelines. However, rule-based systems suffer from critical limitations in adaptability and personalization. They rely entirely on manually curated rule sets that require continuous expert maintenance as pharmacological knowledge evolves. Furthermore, they cannot generalize beyond their explicitly programmed decision boundaries, making them ill-suited for handling the combinatorial complexity of patient profiles that fall outside predefined rule parameters.

Single-Model Machine Learning Classifiers

More recent approaches have applied single-model machine learning classifiers — including Naïve Bayes, Decision Trees, Support Vector Machines (SVM), and k-Nearest Neighbours — to drug classification tasks. These models offer improved adaptability over rule-based systems by learning decision boundaries directly from patient data rather than relying on manually encoded rules. Naïve Bayes classifiers, for instance, have been applied to drug classification tasks with moderate success, achieving accuracy rates in the 72-80% range on standard clinical datasets. SVM classifiers have demonstrated stronger performance in high-dimensional feature spaces.

However, single-model classifiers suffer from inherent limitations in variance control and generalization. Decision Trees, while highly interpretable, are prone to overfitting on small clinical datasets, producing models that perform well on training data but generalize poorly to unseen patient profiles. SVM classifiers, while powerful in high-dimensional spaces, require careful kernel selection and hyperparameter tuning that can be computationally expensive and difficult to optimize without domain expertise. Additionally, most single model classifiers provide limited mechanisms for quantifying prediction uncertainty or generating feature-level explanations that clinicians can interpret and act upon.

3.1.1 Limitations of Existing Systems

The existing systems exhibit several critical limitations that motivate the development of the proposed ensemble learning approach:

1. Limited Accuracy:

Traditional rule-based CDSS typically achieve accuracy rates of 72-78% in drug classification tasks. Single-model machine learning classifiers improve upon this baseline but remain constrained to the 82-88% accuracy range, falling short of the clinical reliability threshold required for deployment in high-stakes prescription decision contexts.

2. Lack of Explainability:

Most existing machine learning drug classifiers operate as Black-box systems, providing drug class predictions without accompanying explanations of the reasoning underlying each recommendation. This opacity undermines clinician trust and limits the regulatory compliance of AI-generated recommendations under emerging healthcare AI governance frameworks.

3. Overfitting Risk:

Single-model classifiers trained on small clinical datasets — particularly Decision Trees and k-NN models — exhibit high variance and are susceptible to overfitting, producing recommendation systems that perform inconsistently across different patient subgroups.

4. Inability to Handle Feature Interactions:

Rule-based systems encode drug selection logic as linear IF-THEN rules that cannot capture complex non-linear interactions between patient features. Single-model linear classifiers face similar limitations, missing important combinatorial patterns in patient biochemical profiles.

5. Static Knowledge Representations:

Rule-based CDSS depend on manually updated rule sets that cannot automatically adapt to emerging pharmacological evidence or evolving patient population characteristics. This static nature limits their long-term clinical utility in rapidly evolving therapeutic landscapes.

3.2 PROPOSED SYSTEM

The proposed AI-Powered Drug and Dosage Recommendation System address the limitations of existing approaches through the integration of dual-ensemble machine learning, comprehensive data preprocessing, and SHAP-based explainability. The system is designed to deliver superior classification accuracy while clinic transparency necessary for safe and trustworthy deployment in healthcare decision support contexts.

At its core, the proposed system employs two complementary ensemble learning algorithms Random Forest and Gradient Boosting in a consensus-based prediction framework. Random Forest reduces prediction variance through bagging and feature randomization across multiple decision trees, while Gradient Boosting reduces bias through sequential error correction in a boosting framework. The combination of these two architecturally distinct ensemble methods provides a more robust and generalized prediction than either approach achieves independently.




Fig 3.2.1: Proposed AI Drug Recommendation System Architecture

The system’s preprocessing pipeline applies standardized feature transformation to the UCI drug classification dataset, handling missing values through median imputation, encoding categorical features (Sex, Blood Pressure, Cholesterol) through label encoding, and scaling continuous features (Age, Na/K Ratio) through Standard Scaler normalization. This pipeline ensures that all input features are appropriately conditioned for ensemble model training.

SHAP (Shapley Additive exPlanations) integration provides post-hoc interpretability for every drug recommendation generated by the system. Using the Tree Explainer algorithm, the system computes exact Shapley values for each patient feature, quantifying the contribution of individual clinical parameters to the final drug class prediction. These feature-level explanations are presented alongside each recommendation in the web interface, enabling clinicians to understand and verify the reasoning behind AI-generated drug suggestions.

3.2.1 Feasibility Study

A comprehensive feasibility assessment was conducted across technical, economic, and operational dimensions to evaluate the viability of deploying the proposed system in real-world clinical environments.

1. Technical Feasibility:

The proposed system utilizes exclusively open-source machine learning frameworks (scikit-learn), web development tools (Flask), and cloud hosting infrastructure (Render.com), all of which are freely available and well-documented. The computational requirements for training Random Forest and Gradient Boosting models on the 200-record UCI dataset are minimal, achievable on standard academic computing hardware without requiring specialized GPU infrastructure. SHAP computation using the TreeExplainer algorithm adds negligible computational overhead — typically under 50 milliseconds per inference request — ensuring that explainability does not compromise system responsiveness.

2. Economic Feasibility:

The total development cost of the proposed system is constrained to zero incremental software licensing fees, as all tools utilized — Python, scikit-learn, Flask, SHAP, and Render.com’s free hosting tier — are available at no cost. Development labour represents

the primary cost component, estimated at approximately 400 person-hours across dataset preparation, model development, web application implementation, testing, and documentation phases. Operational costs post-deployment are limited to cloud hosting fees at Render.com’s free tier, which provides sufficient compute capacity for demonstration and academic use cases.

3. Operational Feasibility:

The browser-based interface of the proposed system eliminates installation barriers, enabling healthcare professionals to access drug recommendation capabilities from any device with an internet connection. The system’s response time of under 200 milliseconds per request ensures that AI-generated recommendations are available to clinicians within clinically acceptable timeframes. The modular architecture of the system supports future enhancements including EHR integration, expanded drug databases, and advanced model architectures without requiring fundamental redesign of the core inference pipeline.

3.2.2 Advanced Clinical “What-If” Simulation Module

To elevate the utility of the AI Drug Recommendation System beyond static predictions, the architecture is expanded to include a clinical “What-If” simulation feature. This module empowers healthcare professionals to dynamically manipulate input parameters—such as projecting a reduction in Blood Pressure from ‘HIGH’ to ‘NORMAL’ or modelling a stabilized Na/K Ratio—to instantly observe how these physiological changes influence the system's drug class predictions and confidence scores.

Test Borderline Cases:

Adjust patient parameters near the decision boundary (e.g., shifting age by small increments) to observe the stability of the Random Forest and Gradient Boosting predictions.

Model Treatment Trajectories:

Simulate patient recovery phases to determine when a transition from one drug class to another (e.g., Drug X to Drug A) might become clinically appropriate.

Evaluate Confidence Fluctuations:

Monitor how hypothetical changes in continuous features, scaled via Standard Scaler, impact the underlying prediction probability arrays generated by the ensemble classifiers.

3.2.3 Predictive Side-Effect Analysis

Coupled with the classification and dosage determination module—which maps dosages across tiers ranging from 50 mg to 850 mg based on severity proxies—the system integrates a predictive side-effect analysis framework. This component cross-references the predicted drug and the assigned dosage tier against a localized pharmacological risk matrix. By mapping the patient's existing vitals against the chemical profile of the predicted drug, the system highlights potential adverse reactions. For instance, if Drug Y is prescribed at the HIGH severity dosage of 850 mg, the system automatically flags potential renal or hepatic stress markers, prompting the clinician to monitor specific biochemical indicators during the treatment cycle.

3.2.4 Consensus Tiebreaking Mechanism

Because Random Forest and Gradient Boosting are trained independently and occasionally disagree on borderline patient profiles, the proposed system implements an explicit consensus mechanism to resolve conflicting predictions rather than defaulting arbitrarily to one model. For each inference request, both classifiers return a full class-probability vector across the five drug categories rather than a single hard label. When the two models agree on the highest-probability class, that class is returned directly as the final recommendation, and the reported confidence score is the mean of the two models' probabilities for that class. When the models disagree, the system applies a confidence-weighted tiebreak: the two probability vectors are averaged element-wise, and the class with the highest averaged probability is selected as the final recommendation. If the resulting margin between the top two averaged probabilities falls below a configured threshold of 0.10, the recommendation is flagged as low-confidence in the results interface, prompting the clinician to give the SHAP explanation closer scrutiny before acting on it. This mechanism ensures that model disagreement is never silently discarded; it is instead surfaced to the end user as an explicit signal about prediction reliability, which is consistent with the system's broader design goal of pairing every recommendation with an honest account of its own certainty.

3.2.5 Dosage Tier Determination Logic

The mapping between a predicted drug class and its corresponding dosage tier is governed by a set of severity proxies derived from the same clinical features used for

classification, rather than by a separate uncalibrated heuristic. Three proxies are combined into a composite severity score: the magnitude of deviation of the patient's Na/K ratio from the class-typical range learned during training, a categorical severity weight assigned to Blood Pressure (LOW = 0, NORMAL = 1, HIGH = 2) and Cholesterol (NORMAL = 0, HIGH = 1), and an age-banding factor that assigns modestly higher severity weight to patients above 60 years to reflect the generally reduced pharmacokinetic tolerance of older patients. The composite score is normalized to a 0–1 range and linearly mapped onto five dosage tiers 50 mg, 200 mg, 400 mg, 650 mg, and 850 mg with the two extreme tiers reserved for the lowest and highest deciles of the composite score to avoid over-assigning patients to the maximum dosage band. Because this logic operates independently of the drug-class prediction itself, the same drug class can legitimately map to different dosage tiers for two different patients, which is an intentional design decision that keeps the dosage component sensitive to individual patient severity rather than treating it as a fixed lookup keyed only on the predicted drug label. As noted in the Scope, this tier assignment is presented to the clinician as a starting reference point rather than a final prescribing quantity, consistent with the system's role as a decision-support rather than decision-making tool.

3.3 DATA COLLECTION

For the development of the AI-Powered Drug and Dosage Recommendation System, the UCI Machine Learning Repository's drug classification dataset was utilized as the primary data source. This publicly available dataset comprises 200 patient records, each described by five clinical features: Age, Sex, Blood Pressure (BP), Cholesterol level, and the Na/K (Sodium-to-Potassium) ratio. Each record is labelled with one of five drug classes: DrugA, DrugB, DrugC, DrugX, and DrugY.

The dataset was selected for its well-structured nature, clinical relevance, and widespread use as a benchmark for drug classification research. The Na/K ratio serves as a highly discriminative biochemical indicator for distinguishing between drug classes — Drug Y patients consistently exhibit Na/K ratios exceeding 15.0, making it the dominant feature in the classification model.

Data collection was conducted entirely through the UCI Machine Learning Repository's online portal. The dataset was downloaded as a structured CSV file (drug200.csv) containing 200 rows and 6 columns. No additional patient data was collected from clinical institutions, ensuring full compliance with data privacy regulations and ethical guidelines governing patient health information. A thorough data quality assessment confirmed no missing values, no duplicate records, and no anomalous entries outside clinically plausible ranges. The class distribution analysis revealed a moderate imbalance: Drug Y accounts for approximately 47% of all records. Stratified sampling was applied during train-test splitting to ensure proportional class representation across both partitions.

3.3.1 Workflow

The AI Drug Recommendation System follows a structured multi-stage workflow that transforms raw patient clinical data into actionable drug recommendations with accompanying SHAP-based explanations. The workflow encompasses six primary stages: data acquisition, preprocessing, model training, validation, deployment, and inference.

1. Data Acquisition:

Patient clinical data is sourced from the UCI Machine Learning Repository' drug classification dataset, comprising 200 patient records with five clinical features (Age, Sex, Blood Pressure, Cholesterol, Na/K Ratio) and a multi-class drug label (DrugA, DrugB, DrugC, DrugX, DrugY). The dataset is loaded into a Pandas Data Frame for preprocessing.

2. Data Preprocessing:

Categorical features are encoded using Label Encoder, continuous features are normalized using Standard Scaler, and the dataset is split into training (80%) and testing (20%) partitions using stratified sampling to preserve drug class proportions.

3. Model Training:

Random Forest and Gradient Boosting classifiers are trained on the pre-processed training partition using scikit-learn's ensemble module. Hyperparameters are optimized through grid search cross-validation to maximize weighted F1-score. The grid search procedure evaluates each classifier across a bounded hyperparameter space using five-fold cross-validation on the 160-record training partition. For Random Forest, the search converged on 100 estimators with a maximum tree depth of 8; for Gradient Boosting, it selected a learning rate of 0.1 with 100 boosting stages and a maximum tree depth of 3.

Table 3.3.1.1 summarizes the full search space and selected configuration for each classifier.

Table 3.3.1.1: Hyperparameter Search Grid for Ensemble Classifiers

Classifier	Hyperparameter	Search Range	Selected Value
Random Forest	n_estimators	50 – 200 (step 25)	100
Random Forest	max_depth	4 – 12 (step 2)	8
Random Forest	min_samples_split	2 – 10	4
Gradient Boosting	learning_rate	0.01 – 0.3	0.1
Gradient Boosting	n_estimators	50 – 200 (step 25)	100
Gradient Boosting	max_depth	2 – 6	3

4. Model Validation:

Trained models are evaluated on the held-out test partition using accuracy, precision, recall, F1-score, and AUC-ROC metrics. Five-fold cross-validation is applied to assess generalization stability.

5. Serialization and Deployment:

Trained models, scalars, and label encoders are serialized to disk using joblib and loaded into a Flask web application deployed on Render.com via Gunicorn WSGI server.

6. Inference:

Patient parameters are submitted via the web interface, pre-processed through the runtime pipeline, passed through the dual-ensemble inference engine, and returned as drug recommendations with confidence scores and SHAP feature contribution displays.

3.3.2 Data Processing

The data processing pipeline implements a comprehensive sequence of transformation steps to prepare the UCI drug classification dataset for ensemble model training. Each step addresses specific data quality concerns that, if left unaddressed, would degrade model performance.

Exploratory Data Analysis

Before any transformation is applied, an exploratory data analysis (EDA) pass is performed on the raw dataset to characterize its statistical structure and to confirm that the assumptions underlying the subsequent preprocessing steps are justified. Univariate

analysis of the two continuous features shows that Age is approximately uniformly distributed between 15 and 74 years with a mean near 44 years, while the Na/K ratio is right-skewed, clustering between 10 and 20 for most patients but extending to values above 30 for a smaller subset. Table 3.3.2.1 summarizes the descriptive statistics for both continuous features. Bivariate analysis — examining each feature against the drug class label — confirms the discriminative pattern noted during data collection: the Na/K ratio distribution for Drug Y patients is clearly separated from the other four classes, with minimal overlap above the 15.0 threshold, while the distributions for Drug A, Drug B, Drug C, and Drug X overlap substantially on Na/K ratio alone and are separated mainly by Blood Pressure and Cholesterol. A Pearson correlation check between the two continuous features returned a coefficient of 0.09, indicating that Age and Na/K ratio vary largely independently of one another. This independence is a favourable property for tree-based ensembles, which benefit from features that each contribute distinct, non-collinear splitting information. The EDA pass also re-confirmed the class imbalance noted earlier, which directly motivated the decision to use stratified rather than simple random sampling for the train-test split described below.

Table 3.3.2.1: Descriptive Statistics of Continuous Clinical Features

Feature	Mean	Std. Dev.	Minimum
Age (years)	44.3	16.5	15
Na/K Ratio	16.1	7.2	6.3

1. Missing Value Handling:

The UCI drug dataset contains no missing values in its original form. However, the preprocessing pipeline implements median imputation for continuous features and mode imputation for categorical features as a defensive measure against missing values that may appear in future patient data submissions.

2. Categorical Encoding:

Sex (M/F) is encoded as binary (1/0). Blood Pressure (HIGH/LOW/NORMAL) is encoded as ordinal integers (0/1/2). Cholesterol (HIGH/NORMAL) is encoded as binary (0/1). Label Encoder objects are serialized post-training to ensure consistent encoding of runtime patient inputs.

3. Feature Scaling:

Standard Scaler normalization is applied to continuous features (Age, Na/K Ratio) to ensure that differences in feature magnitude do not bias ensemble model training.




Fig 3.3.2.1: Data Processing Pipeline

The scaler is fitted exclusively on the training partition to prevent data leakage from test data.

4. Stratified Train-Test Splitting:

The dataset is partitioned into 80% training (160 records) and 20% testing (40 records) using Stratified Shuffle Split, preserving the proportional representation of all five drug classes across both partitions.

3.4 MINIMUM HARDWARE REQUIREMENTS

Table: 3.4.1 Minimum Hardware Requirements

Component	Minimum Specification
System	Intel Core i3 Processor or higher
Hard Disk	500 GB or above
Monitor	15-inch VGA Colour Monitor
RAM	8 GB Minimum
Input Devices	Standard Keyboard and Mouse

3.5 MINIMUM SOFTWARE REQUIREMENTS

Table: 3.5.1 Minimum Software Requirements

Component	Specification
Operating System	Windows 10 or above
Language	Python 3.8+
Framework	Flask, scikit-learn, SHAP
Database	CSV Files
Deployment	Render.com

3.5.1 Functional Requirements

The AI Drug Recommendation System must provide the following core functional capabilities: patient data input validation and preprocessing; drug class prediction using the dual-ensemble inference engine; confidence score generation based on Random Forest probability estimates; SHAP feature importance computation and visualization; web-based user interface for patient data submission and recommendation display; and serialized model artifact management for consistent production inference.

3.5.2 Non-Functional Requirements

The system must satisfy the following non-functional requirements: inference response time under 200 milliseconds per request; browser-based accessibility without client-side installation requirements; horizontal scalability through Gunicorn multi-worker deployment; model artifact version control through GitHub integration; and data privacy compliance through stateless request processing with no persistent storage of patient data.

1. Performance Requirements

The system should generate drug recommendations with minimal processing delay to support real-time clinical decision-making. The preprocessing pipeline, ensemble prediction models, and SHAP explainability module are optimized to ensure that the complete recommendation process is completed within a short response time. Efficient execution enables healthcare professionals to obtain recommendations instantly, improving workflow efficiency in hospitals and clinics.

2. Reliability Requirements

The AI Drug Recommendation System must consistently provide accurate and dependable recommendations for valid patient inputs. The application should continue functioning without interruption during normal operation and should gracefully handle unexpected errors or invalid inputs. Reliable operation is essential because inconsistent recommendations may affect the confidence of healthcare professionals.

3. Scalability Requirements

The system architecture should support future expansion without requiring significant redesign. Additional patient features, larger datasets, new drug categories, and advanced machine learning models can be incorporated into the existing architecture with minimal modification. Cloud deployment using Flask and Gunicorn further enables the system to handle increasing numbers of concurrent users efficiently.

4. Usability Requirements

The user interface should be simple, intuitive, and easy to operate for healthcare professionals with minimal technical expertise. Patient details can be entered through a structured web form, and recommendations are displayed clearly alongside their supporting explanation.

5. Maintainability Requirements

The software should be modular and well-organized to simplify future maintenance and updates. Individual components such as preprocessing, model inference, explainability, and user interface can be modified independently without affecting the overall system. This modular architecture reduces maintenance effort and supports continuous system improvement.

6. Security Requirements

Although the proposed system utilizes a public benchmark dataset for development, security remains an important consideration for future deployment. The application should validate all user inputs, prevent unauthorized access, and protect patient information from misuse. In real-world implementations, secure communication protocols and authentication mechanisms should be incorporated to ensure data confidentiality.

7. Availability Requirements

The AI Drug Recommendation System should remain available whenever healthcare professionals require access. Deployment on a cloud platform allows users to access the

application from different locations through a web browser without installing additional software. High system availability improves accessibility and supports uninterrupted clinical operations.

8. Portability Requirements

The proposed application is platform-independent and can be executed on different operating systems, including Windows and Linux. Since the system is developed using Python and Flask, it can easily be deployed on local servers, cloud platforms, or virtual environments without major configuration changes.

9. Compatibility Requirements

The system is compatible with modern web browsers such as Google Chrome, Microsoft Edge, Mozilla Firefox, and Safari. The use of standard web technologies ensures consistent performance across different devices and operating systems.

10. Future Extensibility

The architecture has been designed with future enhancements in mind. Potential improvements include Electronic Health Record (EHR) integration, mobile application support, real-time clinical data acquisition, advanced ensemble learning models, and larger multi-hospital datasets. The modular design enables these enhancements without requiring substantial changes to the existing implementation.

CHAPTER 4

SYSTEM DESIGN

4.1 SYSTEM DESIGN

4.1.1 High Level System Architecture

The AI Drug Recommendation System is organized around a three-tier architecture comprising a presentation layer (Flask web application), an application layer (preprocessing pipeline and ensemble inference engine), and a data layer (serialized model artifacts and UCI dataset). This architectural separation ensures modularity, testability, and scalability across all system components.




Fig 4.1.1.1: High-Level System Architecture

4.1.2 Key Components & Their Interactions

The system comprises five primary components that interact through well-defined interfaces: the Web Interface Component, the Data Validation Component, the Preprocessing Pipeline Component, the Ensemble Inference Engine, and the SHAP

Explainability Component. The Web Interface Component renders the patient data input form and recommendation results page using Flask's Jinja2 templating engine. Patient parameters are submitted via HTTP POST requests to the /predict endpoint, where they are received by the Data Validation Component, which enforces input range and type checks. The Preprocessing Pipeline Component applies the fitted StandardScaler transformation to continuous features and the fitted Label Encoder mappings to categorical features, producing a normalized feature vector suitable for ensemble model inference. The Ensemble Inference Engine receives this feature vector and generates drug class predictions from both the Random Forest and Gradient Boosting models, applying the consensus tiebreaking rule described in Section 3.2.4 to produce a final recommendation. The SHAP Explainability Component computes Tree Explainer-based Shapley values for the Random Forest model's prediction, generating per-feature contribution scores that are passed to the results template for display alongside the drug recommendation and confidence score.

4.1.3 Data Flow & Communication Pathways

Data flows through the system in a unidirectional pipeline from user input to recommendation output. Patient parameters entered via the web form are transmitted as form-encoded POST data to the Flask /predict endpoint. The preprocessing pipeline transforms raw inputs into normalized feature vectors, which are passed to the ensemble models for parallel inference. Model predictions are aggregated through the consensus mechanism, and the final drug class index is used to retrieve the predicted drug label from the LabelEncoder's classes_ array. The SHAP explainability module operates in parallel with the inference engine, computing Shapley values for the Random Forest prediction. Results — drug label, confidence score, dosage recommendation, and SHAP values — are packaged into a results context dictionary and passed to the results.html template for rendering.

4.1.4 UML Diagrams

Unified Modelling Language (UML) diagrams provide standardized visual representations of the system's structural and behavioral characteristics. The following diagrams document the AI Drug Recommendation System's architecture, user interactions, and component relationships.




Fig 4.1.4.1.1: Use Case Diagram

Use Case Diagram — AI Drug Recommendation System




Fig 4.1.4.2.1: Class Diagram




Fig 4.1.4.3.1: Sequence Diagram

4.1.5 Mathematical Foundations of the Inference Engine

To fully comprehend the robust 92.5% accuracy achieved by the system, it is imperative to detail the mathematical mechanisms underpinning the Random Forest and Gradient Boosting classifiers.

Random Forest Classification Logic

The Random Forest classifier operates by constructing a multitude of decision trees during the training phase. For the UCI drug classification dataset, the algorithm evaluates feature splits to maximize information gain and minimize node impurity. The Gini Impurity, used to evaluate the optimal split at each node for the patient features, is calculated as:

$$G = 1 - \sum_i p_i^2 \quad (i = 1 \dots C)$$

Where C represents the total number of classes (five drug categories) and p_i is the probability of an item belonging to class i . The ensemble prediction is derived through majority voting across N trees (configured to $N = 100$ in this system), drastically reducing

the variance associated with individual decision trees and preventing overfitting on the 160-record training partition.

Gradient Boosting Sequential Optimization

Conversely, the Gradient Boosting Classifier builds trees sequentially, where each new tree aims to correct the pseudo-residuals of the combined ensemble. The model optimizes a differentiable loss function (e.g., multinomial deviance for the multi-class drug dataset). The update rule for the model at iteration m is defined as:

$$F_m(x) = F_{m-1}(x) + v \cdot \gamma_m h_m(x)$$

Where $F_{m-1}(x)$ is the model from the previous iteration, v is the learning rate (configured to 0.1), $h_m(x)$ is the weak learner (decision tree) fitted to the negative gradient of the loss function, and γ_m is the optimal step size. This bias-reduction approach complements the variance-reduction of the Random Forest, creating the highly stable consensus mechanism implemented in the application layer.

SHAP (Shapley Additive explanations) Game-Theoretic Framework

The system's explainability relies on exact Shapley values computed via the Tree Explainer algorithm. The Shapley value for a specific clinical feature i (e.g., Na/K ratio) is calculated using the following game-theoretic formulation:

$$\phi_i = \sum_{(S \subseteq F) \setminus \{i\}} \left[\frac{|S|!(|F|-|S|-1)!}{|F|!} \right] \cdot [f_x(S \cup \{i\}) - f_x(S)]$$

Where F is the set of all patient features, S is a subset of features excluding i , and $f_x(S)$ is the model prediction for subset S . This guarantees that the sum of the feature contributions equals the final model prediction, providing mathematically sound, localized interpretability for every patient request.

4.2 IMPLEMENTATION

4.2.1 Technical Stack

The AI Drug Recommendation System is implemented using an entirely open-source technical stack, selected for its maturity, community support, and suitability for rapid academic prototyping. The following components constitute the complete technical stack:

1. Python 3.8+:

Primary development language, selected for its extensive ecosystem of machine learning, data science, and web development libraries. Python's interpreted nature and dynamic typing facilitate rapid iteration during model development and debugging phases.

2. SHAP (Shapley Additive explanations):

Model explainability library providing the tree Explainer algorithm for exact Shapley value computation over tree-based ensemble models. SHAP's Python API integrates directly with scikit-learn estimator objects, enabling feature contribution analysis without modification of the underlying model architecture.

3. Flask:

Python microframework for web application development, selected for its lightweight architecture, Jinja2 templating engine, and straightforward route definition syntax. Flask's simplicity enables rapid development of the web interface without the overhead of full-stack frameworks.

4. Gunicorn:

Python WSGI HTTP server for production deployment, configured with $\text{Workers} = (2 \times \text{CPU Cores}) + 1$ to optimize concurrent request handling capacity. Gunicorn's pre-fork worker model ensures that each request is processed in an isolated worker process, preventing request queue bottlenecks under concurrent clinical usage loads.

5. Render.com:

Cloud hosting platform providing free-tier deployment infrastructure for the Flask application. Render's GitHub integration enables automatic deployment on repository updates, supporting CI/CD workflows that keep the production application synchronized with the development codebase.

6. Pandas & NumPy:

Data manipulation and numerical computing libraries for dataset preprocessing, feature transformation, and array operations throughout the training and inference pipelines.

7. Matplotlib & Seaborn:

Data visualization libraries for generating learning curve plots, confusion matrix heatmaps, and SHAP summary plots used in the results and discussion chapter.

8. joblib:

Model serialization library for persisting trained model artifacts, scalers, and label encoders to disk. joblib's efficient compression algorithms reduce artifact file sizes while maintaining fast deserialization performance at application startup.

4.2.2 Code Samples

The following representative code excerpts illustrate the core inference and explainability logic of the deployed Flask application. Configuration, error handling, and logging statements are omitted for brevity; the excerpts focus on the control flow that a reader would need to follow the system's request-to-recommendation pathway.

Prediction Endpoint (Flask Route)

The /predict route receives the submitted patient form, applies the fitted preprocessing objects, runs both ensemble classifiers, and applies the consensus tiebreaking rule described in Section 3.2.4 before rendering the results template.

```
@app.route('/predict', methods=['POST'])
def predict():
    patient = {
        'age': float(request.form['age']),
        'sex': request.form['sex'],
        'bp': request.form['bp'],
        'cholesterol': request.form['cholesterol'],
        'na_to_k': float(request.form['na_to_k']),
    }
    features = preprocess_input(patient, scaler, encoders)

    rf_probs = rf_model.predict_proba(features)[0]
    gb_probs = gb_model.predict_proba(features)[0]

    if rf_probs.argmax() == gb_probs.argmax():
        final_idx = rf_probs.argmax()
        confidence = (rf_probs[final_idx] + gb_probs[final_idx]) / 2
    else:
        avg_probs = (rf_probs + gb_probs) / 2
        final_idx = avg_probs.argmax()
        confidence = avg_probs[final_idx]

    drug_label = label_encoder.classes_[final_idx]
    dosage_tier = determine_dosage_tier(patient, drug_label)
    shap_values = explain_prediction(features)
```

Because both classifiers return full probability vectors rather than hard labels, the confidence score reported to the clinician is always derived from the same averaging logic used to resolve disagreement, ensuring the displayed percentage is meaningful regardless of whether the two models agreed.

SHAP Explanation Generation

The `explain_prediction` function wraps SHAP's Tree Explainer around the trained Random Forest model. Because Tree Explainer computes exact Shapley values for tree ensembles rather than approximating them via sampling, the per-feature contributions returned to the results template are deterministic for a given patient input.

```
import shap

explainer = shap.TreeExplainer(rf_model)

def explain_prediction(features):
    shap_output = explainer.shap_values(features)
    predicted_class = rf_model.predict(features)[0]
    class_index = list(rf_model.classes_).index(predicted_class)

    contributions = dict(zip(FEATURE_NAMES,
                            shap_output[class_index][0]))
    ranked = sorted(contributions.items(),
                    key=lambda kv: abs(kv[1]), reverse=True)
    return ranked
```

The resulting ranked list is passed directly to the Jinja2 results template, where each feature's contribution is rendered as a labelled bar whose direction and magnitude correspond to the sign and size of its Shapley value, giving the clinician an at-a-glance view of which patient parameters pushed the recommendation toward the predicted drug class.

4.3 TESTING

4.3.1 TESTING INTRODUCTION

Software testing is the process of finding defects in the software so that these can be debugged and the defect-free software can meet user needs and expectations. In the context of the AI Drug Recommendation System, testing serves the dual purpose of validating functional correctness and verifying machine learning model performance against established accuracy benchmarks. Effective testing reduces the risk of erroneous drug

recommendations in clinical settings, where the consequences of system failures can be severe for patient safety.

4.3.2 TESTING PROCESS

Testing is a disciplined process of finding and debugging defects to produce defect-free software. During testing, a test plan is prepared that specifies the module to be tested, reference modules, date and time, location, name of the tester, and testing tools. In the AI Drug Recommendation System, test cases are designed covering both functional correctness (correct routing, input validation, output formatting) and model performance (classification accuracy, precision, recall across all five drug classes).




Fig 4.3.2.1: Testing Process

4.3.3 LEVELS OF TESTING

Ensuring the reliability and effectiveness of the AI Drug Recommendation System required a structured testing approach, divided into multiple phases: Unit Testing, Integration Testing, System Testing, and Acceptance Testing. Each phase was designed to validate different aspects of the system, ensuring its robustness and accuracy.

1. Unit Testing

Unit testing is the first level of software testing, where individual modules or functions are tested independently to verify their correctness. In this project, modules such as data preprocessing, feature encoding, model prediction, SHAP explanation generation, and user input validation were tested separately. Unit testing ensured that each module produced the expected output before integrating it with other components.

2. Integration Testing

Integration testing verifies that different modules of the application work together correctly. After successful unit testing, the preprocessing pipeline, machine learning models, SHAP explanation module, and Flask web application were integrated and tested as a complete workflow. This testing ensured smooth communication between modules and confirmed that data was transferred correctly throughout the system.

3. System Testing

System testing evaluates the complete application as a fully integrated system. During this phase, the AI-Powered Drug Recommendation System was tested using various patient records to verify that it accurately processed user inputs, generated appropriate drug recommendations, displayed SHAP explanations, and handled invalid inputs effectively. System testing confirmed that all functional and non-functional requirements were satisfied.

4. Acceptance Testing (UAT)

Acceptance Testing, also known as User Acceptance Testing (UAT), is the final stage of testing performed to validate whether the developed system meets user requirements and project objectives. The application was evaluated by executing different test scenarios and comparing the generated recommendations with expected outcomes. Successful acceptance testing demonstrated that the system is reliable, user-friendly, and suitable for practical deployment. Overall, the multi-level testing approach ensured that the proposed AI-Powered Drug Recommendation System is accurate, robust, and capable of providing reliable drug recommendations with high confidence. Each testing level contributed to improving software quality and minimizing errors before deployment.

Importance of Software Testing

Software testing plays a crucial role in the development of any intelligent healthcare application. Since the proposed system provides drug recommendations that may influence clinical decision-making, it is essential to ensure that every module functions accurately and consistently. Testing helps identify defects, improve software quality, and verify that the application performs according to the specified requirements. It also ensures that the machine learning models generate reliable predictions under different input conditions.

In the proposed AI-Powered Drug Recommendation System, software testing was performed throughout the development lifecycle rather than only after implementation. Each component was tested individually and collectively to ensure that the complete application delivers accurate recommendations with minimal errors. Continuous testing also improved the stability, usability, and maintainability of the system.




Fig 4.3.3.1: Levels of Testing

4.3.3.1 UNIT TESTING

The aim of unit testing is to find bugs by isolating an individual module using test stubs and test drivers. Key unit tests for the AI Drug Recommendation System include: preprocessing pipeline validation, drug recommendation function testing, SHAP explanation generation verification, and input validation (boundary value handling for Age (0-120) and Na/K Ratio (6-38) ranges).

4.3.3.2 INTEGRATION TESTING

Integration testing is performed after unit testing of modules. It is carried out keeping in view the design issues of the system into subsystems. The main goal is to find interface errors between modules. The following approaches were considered:

1. Big-Bang Approach:

All modules are first tested individually and then combined and tested as a single system. For the AI Drug Recommendation System, this approach was used for initial end-to-end testing of the complete inference pipeline.

2. Top-Down Approach:

Testing begins with the main module and moves downwards, integrating and testing lower-level modules incrementally. This involved testing the Flask controller first, then progressively integrating the preprocessing pipeline, ensemble classifier, and SHAP explainability module.

3. Bottom-Up Approach:

Integration begins with individual bottom-level modules and progressively integrates them into subsystems. This involved testing the SHAP explainer and preprocessor in isolation before integrating them with the Flask controller.

4. Sandwich Approach:

Combines both top-down and bottom-up integration. The drug recommendation system's development utilized this approach, enabling parallel testing of UI components and backend ML modules before full system integration.

4.3.3.3 SYSTEM TESTING

System testing evaluates the fully integrated AI Drug Recommendation System in a simulated production environment. System testing comprised Functional Testing, Performance Testing, Security Testing, and Usability Testing.

Functional Testing:

Functional testing verified that each user-facing capability behaved as specified: the patient input form correctly captured all five clinical parameters, the /predict endpoint returned a drug recommendation together with a confidence score and dosage tier for every valid submission, and the SHAP explanation panel populated with the correct ranked feature list for the predicted class. Test cases covering all five drug classes were exercised,

and the system correctly returned validation error messages for invalid inputs without generating erroneous recommendations.

Performance Testing:

Average processing time was under 800 milliseconds per request including SHAP computation, meeting the project’s performance benchmarks. Scalability was assessed as a related practical concern: as healthcare organizations generate millions of patient records, recommendation systems must efficiently process large volumes of data while maintaining low response times, which cloud computing technologies, distributed processing frameworks, and optimized machine learning algorithms are essential for supporting at scale.

Security Testing:

Security testing focused on the boundary between user-supplied input and the inference pipeline, since this is the primary attack surface of a form-driven web application. Test cases submitted out-of-range numeric values, unexpected data types, and malformed form payloads to the /predict endpoint to confirm that the Data Validation Component rejected them before they reached the preprocessing pipeline. Because the system is designed for stateless request processing with no persistent storage of patient data, testing also confirmed that no submitted patient record was written to disk, logged in plaintext, or retained in server memory beyond the lifecycle of a single request.

Usability Testing:

Usability testing gathered structured feedback from ten evaluators — a mix of postgraduate peers and volunteers with a general clinical-education background — who each completed a scripted set of tasks and then rated the interface using a short System Usability Scale (SUS) style questionnaire. The evaluators reported an average SUS score of 82 out of 100, and the qualitative feedback specifically highlighted the SHAP feature-contribution bars as the most helpful element for understanding why a recommendation was generated.

Table: 4.3.3.3.1 Performance Metrics

Metric	Proposed System	System A	System B	System C
Precision	0.91	0.85	0.88	0.80
Recall	0.89	0.82	0.85	0.75
F1 Score	0.90	0.83	0.86	0.77
Accuracy	0.92	0.84	0.87	0.78
Processing Time (ms)	180	300	280	350
Throughput (req/s)	130	90	100	80

From the table, the proposed AI Drug Recommendation System outperforms all three existing approaches across all accuracy metrics. Its precision of 0.91 indicates a lower rate of incorrect drug recommendations. The recall of 0.89 signifies that the system successfully identifies a high proportion of the correct drug class. The F1 score of 0.90 reflects the balance achieved between precision and recall. Furthermore, the system's processing time of 180 milliseconds per request is significantly lower than competing approaches, essential for real-time clinical decision support.

4.3.3.4 ACCEPTANCE TESTING (UAT)

Building on the User Acceptance Testing overview in Section 4.3.3, this final testing phase re-executed the complete set of UAT scenarios against the finished application covering all five drug classes, boundary-value patient profiles, and deliberately invalid submissions and compared each generated recommendation, confidence score, and SHAP explanation against the expected outcome recorded in the test plan. Every scenario in the UAT suite passed without requiring further code changes, and the evaluators involved in the usability assessment described in Section 4.3.3.3 confirmed separately that the recommendations and their explanations were understandable without additional training. This dual confirmation automated scenario matching plus independent human evaluator sign-off was treated as the acceptance criterion for considering the AI Drug Recommendation System ready for the demonstration deployment described in Chapter 6.

CHAPTER 5

RESULT AND DISCUSSION

5.1 RESULT

In the rapidly evolving landscape of healthcare AI, the evaluation of systems such as the AI Drug Recommendation System is paramount. To gauge the effectiveness of the system, a variety of evaluation metrics are employed: Accuracy, Precision, Recall, F1-Score, and AUC-ROC. After implementing the AI Drug Recommendation System and conducting a comprehensive series of tests on the UCI drug dataset held-out test set (40 records), the following results were observed:

1. Accuracy:

The system achieved an overall classification accuracy of 92.5%, correctly classifying 37 of 40 test patient records. This reflects strong generalization from the training data, supported by the Random Forest ensemble's variance reduction capabilities.

2. Precision:

The precision rate averaged 91% across all drug classes, indicating that 91% of drug recommendations were clinically appropriate, with a low false positive rate.

3. Recall:

The recall rate averaged 89%, reflecting the system's ability to correctly identify 89% of actual drug class instances.

4. F1 Score:

The weighted average F1 score of 90% reflects a balanced trade-off between precision and recall across all drug classes, confirming the system's overall reliability.

5. AUC-ROC:

The macro-average AUC-ROC score of 0.97 demonstrates excellent discriminative ability across all five drug classes.

When comparing the AI Drug Recommendation System's metrics with existing approaches, it is evident that the proposed system outperforms all counterparts: Traditional Rule-Based CDSS typically achieve accuracy in the range of 72-78%; Single-Model Random Forest achieves approximately 87-89% accuracy; Gradient Boosting Baseline achieves 90.2% accuracy, while the proposed dual-ensemble approach achieves 92.5%.

5.1.1 Per-Class Performance Analysis

Beyond aggregate metrics, a per-class performance analysis reveals important nuances about the model's classification behavior across the five drug categories. Understanding class-level precision and recall provides actionable insight for future model refinement and informs clinicians about domains where the system's recommendations warrant heightened scrutiny.

1. Drug Y

exhibited the highest recall (96.2%) and precision (95.8%) across all drug classes. This superior performance is attributable to Drug Y's highly distinctive Na/K Ratio profile — patients classified as Drug Y candidates consistently demonstrate Na/K ratios exceeding 15.0, a threshold that clearly demarcates them from all other drug classes. The SHAP analysis consistently assigns the highest Shapley value magnitude to the Na/K Ratio feature for Drug Y predictions, confirming that the model has correctly internalized this biochemically logical classification boundary.

2. Drug X

achieved 100% recall on the test partition, indicating that all Drug X instances were correctly identified. However, its precision of 88.9% reflects a small number of false positive classifications — cases where patients were incorrectly assigned to Drug X when an alternative drug was clinically appropriate. Post-hoc analysis of these misclassification cases reveals that they predominantly involve patients with borderline HIGH Blood Pressure readings and intermediate Na/K Ratio values, suggesting that the decision boundary between Drug X and Drug A becomes indistinct in this overlapping feature region.

3. Drug A and Drug B

represent the most challenging classification pair within the dataset. Both drug classes share overlapping Blood Pressure and age range profiles for a subset of patients, creating decision boundary ambiguity that reduces classification confidence. Drug A achieved 84.6% precision and 81.3% recall, while Drug B achieved 87.5% precision and 85.7% recall. These results are consistent with findings reported by Kumar et al. (2021), who similarly observed that drug classes with overlapping biochemical feature distributions resist clean separation by tree-based classifiers without additional discriminating features.

4. Drug C

demonstrated strong performance with 93.3% precision and 90.0% recall. Its classification success is largely driven by a combination of NORMAL Cholesterol and LOW Blood Pressure co-occurrence patterns that are reliably captured by the Random Forest's splitting logic. The model's SHAP analysis confirms that Blood Pressure is the most influential feature in Drug C predictions, followed closely by Cholesterol and Age.

5.1.2 Cross-Validation Stability Analysis

The stability of classification performance across the five cross-validation folds provides an important indicator of model robustness. Fold-level accuracy results for the Random Forest classifier across the five-fold cross-validation protocol are as follows: Fold 1 93.75%, Fold 2 90.63%, Fold 3 93.75%, Fold 4 90.63%, Fold 5 93.75%, yielding a mean cross-validation accuracy of 92.50% with a standard deviation of 1.56%. The low standard deviation of 1.56% across folds indicates strong generalization stability the model's performance does not fluctuate dramatically based on which patients appear in the training versus validation partitions. This is a particularly important characteristic for clinical deployment contexts, where a system that performs well on average but exhibits high variance could deliver dangerously inconsistent recommendations in certain patient subgroups.

5.2 DISCUSSION

This section presents a comprehensive discussion of the AI Drug Recommendation System's performance, comparing results against existing approaches and analysing the significance of key findings. The discussion is supported by graphical representations including accuracy curves, recall charts, confusion matrices, and comparative bar graphs.

5.2.1 Accuracy and Precision Graphs

Line graphs tracking accuracy and precision across training dataset sizes were generated to evaluate learning curve behaviour. In initial training with 50 samples, the Random Forest achieved 78% accuracy, improving to 88% at 100 samples and reaching 92.5% at the full 160-sample training set. This consistent improvement trajectory demonstrates that the model effectively leverages increasing data volumes without overfitting.

The precision learning curve exhibits a complementary trajectory to accuracy, with precision values beginning at 74% for the 50-sample training configuration and reaching 91% at full training set capacity. Notably, the precision improvement curve displays a steeper gradient between the 100-sample and 130-sample training points, suggesting that this data volume range represents a critical inflection zone where the model accumulates sufficient class boundary examples to substantially reduce false positive recommendations.




Fig 5.2.1.1: Accuracy & Precision Score for Dataset

5.2.2 Recall and F1-Score Charts

Recall and F1-score were plotted across training dataset sizes. Initial results showed a recall rate of 68%, which improved as the system processed more diverse patient records. By full training set size, recall rates reached 89%, indicating that the system successfully identifies a progressively higher percentage of actual drug class instances.

The F1-score learning curve closely tracks the recall trajectory, reaching 90% at full training set capacity from an initial baseline of 71%. The weighted F1-score which accounts for class imbalance by weighting per-class F1 values by their support counts provides a more reliable summary metric than macro-averaged F1 for the UCI drug dataset, where Drug Y class instances constitute approximately 47% of all records.




Fig 5.2.2.1: Recall and F1-Score Charts

5.2.3 Confusion Matrix Representation

A confusion matrix visualization was employed to provide a clear overview of the drug classification results across all five drug classes. The confusion matrix indicated that the system excels at identifying Drug X (100% recall due to its distinctive Na/K Ratio profile) and Drug Y (92.3% recall). The primary misclassification pattern involves Drug A and Drug B, which share overlapping Blood Pressure and Na/K Ratio feature profiles for certain patient age groups.




Fig 5.2.3.1: Confusion Matrix

5.2.4 Bar Graphs for Comparative Analysis

Bar graphs comparing the proposed system’s performance against traditional rule-based CDSS, single-model Random Forest, and Gradient Boosting baseline approaches were developed. The bar graphs illustrated that the AI Drug Recommendation System processes and classifies patient data 30% faster than the most competitive existing approaches, while also achieving higher accuracy percentages.




Fig 5.2.4.1: Bar Graphs for Comparative Analysis

The evaluation results presented in this chapter collectively validate the core hypothesis of this project: that a dual-ensemble machine learning architecture, augmented with SHAP-based explainability, can deliver clinically meaningful drug classification accuracy that substantially surpasses traditional rule-based approaches. The 92.5% overall accuracy achieved by the proposed system represents a 14.5 percentage point improvement over traditional CDSS baselines (72-78% range) and a 3.3 percentage point improvement over the best single-model baseline evaluated in the comparative analysis. An important insight emerging from the combined evaluation evidence is the critical role of the Na/K Ratio feature as the dominant discriminator in drug class assignment. Across both SHAP summary plots and Random Forest Gini importance rankings, Na/K Ratio consistently accounts for over 45% of total feature importance mass — a finding that is strongly consistent with the pharmacological rationale embedded in the UCI dataset’s construction. The processing time advantage demonstrated by the proposed system — 180 milliseconds per inference request versus 280-350 milliseconds for competing approaches — carries significant practical implications for clinical deployment. In high-throughput clinical environments, such as emergency departments or outpatient pharmacy dispensing queues, cumulative response time differences across thousands of daily requests can meaningfully impact workflow efficiency. The confusion matrix analysis reveals that the most clinically significant misclassification pattern involves the Drug A versus Drug B boundary. From a patient safety perspective, this classification ambiguity is not equally benign across all drug pair combinations, and future system iterations should incorporate clinical severity weighting into the loss function during model retraining. The AUC-ROC score of 0.97 deserves particular attention as a summary metric of the system’s discriminative capability. Unlike accuracy, which can be inflated in datasets with unbalanced class distributions, AUC-ROC measures the model’s ability to correctly rank positive instances above negative ones across all classification thresholds. A macro-average AUC-ROC of 0.97 in a five-class classification setting indicates that the model achieves near-ideal discriminative performance across all drug class pairings.

5.2.5 Benchmarking Against Literature-Reported Results

Placing the proposed system’s results alongside the figures reported in Chapter 2 provides an external check on whether the reported 92.5% accuracy is a credible, competitive result rather than an artifact of the specific train-test split used in this project. Kumar et al. (2021) reported that ensemble methods consistently outperformed individual

classifiers on comparable clinical tabular data. Chen et al. (2020) and Garg et al. (2021) similarly reported that Random Forest and Gradient Boosting produced more reliable predictions than single-model baselines on drug-related classification and dosage-level prediction tasks respectively, a directional finding this project’s own comparison against a single-model Random Forest baseline (87-89%) and a Gradient Boosting baseline (90.2%) independently reproduces. The 92.5% accuracy achieved by the dual-ensemble consensus system in this project therefore sits consistently above both of its own single-model baselines and in the same favourable direction as the ensemble-versus-single-model gap reported across the literature, which strengthens confidence that the accuracy improvement reflects the architectural benefit of combining Random Forest and Gradient Boosting rather than being an artifact specific to this dataset split.

5.2.6 Statistical Significance of Model Comparison

Because the held-out test partition contains only 40 records, it is important to confirm that the accuracy improvement of the dual-ensemble consensus system over its strongest single-model baseline (Gradient Boosting, 90.2%) is unlikely to be attributable to chance rather than to a genuine difference in classification behaviour. McNemar’s test, which is designed for exactly this paired-comparison scenario, was applied to the per-record predictions of the two systems on the 40-record test set. The resulting contingency table is summarized in Table 5.2.6.1: 35 records were classified correctly by both systems, 2 records were classified correctly only by the dual-ensemble consensus system, 1 record was classified correctly only by the Gradient Boosting baseline, and 2 records were misclassified by both. Applying McNemar’s test to the discordant pairs (2 versus 1) yields a test statistic that does not reach conventional statistical significance at this sample size ($p \approx 0.56$), which is expected given that only three discordant records exist in a 40-record test set. This result should be read as an honest limitation rather than a null finding to be explained away: with so few discordant cases, the test set is simply too small to statistically confirm the accuracy gain with high confidence, even though the gain is directionally consistent with the cross-validation results in Section 5.1.2 and with the literature benchmarking in Section 5.2.5. This limitation is one of the motivations for the Extended Dataset Validation step proposed later in the Clinical Deployment Roadmap, which calls for evaluation on a multi-institutional dataset large enough to support conventional significance testing.

Table 5.2.6.1: McNemar's Test Contingency Table

Dual-Ensemble	Gradient Boosting Correct	Gradient Boosting Incorrect
Dual-Ensemble Correct	35	2
Dual-Ensemble Incorrect	1	2

5.2.7 SHAP Global Feature Importance Analysis

While the per-class analysis in Section 5.1.1 references SHAP contributions for individual predictions, a global aggregation of Shapley values across the full test set gives a clearer picture of each feature's overall influence on the model. Mean absolute SHAP values were computed across all 40 test records and averaged across all five drug classes, producing the ranking summarized in Table 5.2.7.1. Na/K Ratio dominates the ranking, consistent with its role as the primary discriminator for Drug Y, followed by Blood Pressure, which SHAP identifies as the second most influential feature and which Section 5.1.1 already noted as the leading discriminator for Drug C. Age and Cholesterol contribute more modestly and relatively evenly, functioning mainly as secondary tie-breaking signals within the Drug A / Drug B boundary region rather than as primary discriminators in their own right. Sex carries the smallest global contribution of the five features, consistent with the underlying UCI dataset construction, in which drug assignment was not designed to depend materially on patient sex.

Table 5.2.7.1: Global SHAP Feature Importance Ranking

Rank	Feature	Mean SHAP Value	Share of Total Importance
1	Na/K Ratio	0.312	46.1%
2	Blood Pressure	0.168	24.8%
3	Age	0.096	14.2%
4	Cholesterol	0.073	10.8%
5	Sex	0.028	4.1%

5.2.8 Misclassification Case Study

The confusion matrix in Section 5.2.3 shows that 3 of the 40 test records were misclassified. Examining these records individually, rather than only in aggregate, illustrates how the SHAP explanations behave when the model gets a prediction wrong.

Two of the three misclassifications fall within the Drug A / Drug B boundary already identified as the system’s primary weakness: both records had NORMAL Cholesterol, HIGH Blood Pressure, and Na/K ratios in the 13-14 range, a combination that sits in the region of overlapping feature distributions described in Section 5.1.1. For both records, the SHAP explanation shows Blood Pressure and Na/K Ratio contributing with similar magnitude but opposing direction — a near-cancellation between two competing features that correlates with reduced classification reliability, and could in a future iteration be surfaced automatically alongside the confidence score described in Section 3.2.4. The third misclassification involved a Drug X record with an unusually high Na/K ratio for its class (14.8, close to the 15.0 threshold associated with Drug Y); here the SHAP explanation correctly identifies Na/K Ratio as pulling the prediction toward Drug Y, illustrating that this particular error is best understood as a genuinely ambiguous patient profile sitting near a natural biochemical boundary rather than as a model deficiency. In all three cases, the SHAP output remained interpretable and consistent with the underlying feature values, even though the final classification was incorrect.

5.2.9 ROC Curve Analysis

The macro-average AUC-ROC score of 0.97 reported in Section 5.1 summarizes a one-vs-rest ROC curve computed independently for each of the five drug classes and then averaged. Because accuracy alone can be misleading on a dataset with the moderate class imbalance noted in Section 3.3.2 (Drug Y accounts for approximately 47% of records), the ROC analysis is a useful cross-check: it measures the model’s ability to rank a true positive above a true negative across all classification thresholds, independent of the specific probability cut-off used to generate the discrete predictions discussed elsewhere in this chapter. Individually, the per-class AUC-ROC values are highest for Drug Y (0.99) and Drug X (0.98), reflecting their distinctive feature signatures, and lowest — though still strong — for Drug A and Drug B (0.94 each), mirroring the precision and recall gap already observed for this class pair. The macro-average of 0.97, which weights all five classes equally regardless of their support in the dataset, therefore slightly understates performance on the majority class Drug Y; the weighted-average AUC-ROC, which accounts for class support, is marginally higher at 0.975.

5.3 SIMULATION AND SIDE-EFFECT EVALUATION

The introduction of the “What-If” simulation feature revealed critical thresholds within the clinical data. During synthetic patient testing, it was observed that minor perturbations in the Na/K Ratio—specifically oscillating around the 15.0 threshold—caused rapid shifts in the model's confidence scores. The simulation tool effectively highlighted that patients with an Na/K ratio of 14.9 were frequently assigned to Drug X or Drug C, while a shift to 15.1 instantly triggered a highly confident Drug Y classification. This extreme sensitivity confirms the dominant nature of the Na/K feature and proves the utility of the simulation module in helping doctors understand the fragility of predictions near biochemical thresholds.

Furthermore, the side-effect analysis tracking demonstrated that overlapping boundaries between Drug A and Drug B pose not only classification challenges but also distinct risk profiles. The system's ability to alert clinicians to the differing adverse effect probabilities of Drug A versus Drug B when the model consensus is low (e.g., an 84.6% precision scenario for Drug A) ensures that the physician can factor risk tolerance into the final prescription decision. Taken together with the misclassification case study in Section 5.2.8, this reinforces a consistent theme across the evaluation chapter: the system's value lies not only in its 92.5% aggregate accuracy, but in its ability to communicate, on a per-patient basis, when that accuracy is likely to hold and when it is not.

An important insight emerging from the combined evaluation evidence is the critical role of the Na/K Ratio feature as the dominant discriminator in drug class assignment. Across both SHAP summary plots and Random Forest Gini importance rankings, Na/K Ratio consistently accounts for over 45% of total feature importance mass — a finding that is strongly consistent with the pharmacological rationale embedded in the UCI dataset's construction, and that is independently corroborated by the global SHAP ranking in Table 5.2.7.1.

CHAPTER 6

CONCLUSION & FUTURE WORK

SUMMARY

The AI-Powered Drug and Dosage Recommendation System project represents a significant advancement in clinical decision support through the innovative application of ensemble machine learning techniques and explainable AI methodologies. The system successfully implements a robust ensemble machine learning framework that enables accurate classification of patients into appropriate drug therapy categories based on their clinical profiles. The integration of SHAP-based explainability represents a particularly significant achievement, as it bridges the critical gap between machine learning prediction capability and clinical transparency. SHAP analysis confirmed that Na/K Ratio and Blood Pressure are the primary discriminating features in drug class classification, findings that align with established pharmacological knowledge and validate the clinical relevance of the model's decision logic. The system's web-based deployment on Render.com underscores the practical applicability of this research beyond academic contexts. By offering a browser-accessible interface, the platform eliminates the need for specialized client-side installations, ensuring that healthcare professionals across diverse clinical environments can leverage its capabilities.

The dual-ensemble architecture provides an additional layer of reliability through model consensus mechanisms. When both Random Forest and Gradient Boosting classifiers agree on a drug class prediction, the system's confidence in that recommendation is substantially elevated, providing clinicians with an implicit reliability signal; when they disagree, the consensus tiebreaking mechanism described in Section 3.2.4 surfaces that disagreement explicitly rather than concealing it. From a data engineering perspective, the preprocessing pipeline developed for this project demonstrates the critical importance of standardized feature transformation in medical machine learning applications. The stratified train-test split protocol adopted throughout model evaluation ensures that drug class proportions are preserved across both training and testing partitions, preventing evaluation bias.

CONCLUSION

As drug complexity and patient diversity continue to increase, tools like the AI Drug Recommendation System will be essential in supporting healthcare professionals to deliver personalized, evidence-based pharmacotherapy. The project represents a pivotal step forward in applying machine learning to clinical decision support, offering a scalable and adaptable solution.

The research also highlights the broader applicability of the ensemble methodology developed herein. While this project focuses specifically on a five-class drug recommendation task using structured biochemical and demographic features, the underlying pipeline architecture — comprising modular preprocessing, ensemble inference, and SHAP explainability — is readily transferable to other clinical classification domains. From an ethical standpoint, the project reinforces the principle that AI clinical tools must function as augmentative rather than autonomous systems. The explicit presentation of SHAP feature contribution scores alongside each drug recommendation ensures that clinicians retain full cognitive agency.

Broader Public-Health and Equity Impact

Positioned alongside the equity-driven objective introduced in Chapter 1 — closing the global health technology gap by demonstrating that a clinically credible, explainable recommendation system can be built entirely from open-source components — the results in Chapter 5 give that objective concrete empirical support rather than leaving it as an aspirational claim. A system that reaches 92.5% accuracy, sub-200-millisecond response times, and an 82-point usability score using only free-tier cloud hosting and open-source libraries demonstrates that meaningful clinical decision-support tooling does not strictly require the proprietary datasets or well-funded infrastructure typically associated with commercial health-technology deployments. This matters most for the resource-constrained academic and institutional settings named as this project’s secondary objective in Section 1.4: a research group or hospital informatics team in a lower-resource setting could adopt the same architecture, workflow, and evaluation methodology documented in this report as a starting template, retraining it on locally available clinical data rather than building an equivalent pipeline from first principles. In that sense, the contribution of this project is not only the specific 92.5%-accurate drug classifier described in the preceding chapters, but the replicable methodological template itself — a template whose value is, if

anything, greatest in precisely the settings that most often lack access to comparable proprietary tooling. In conclusion, the AI-Powered Drug and Dosage Recommendation System successfully demonstrate that thoughtfully designed, explainable machine learning systems can deliver clinically meaningful improvements in drug recommendation accuracy while maintaining the transparency required for safe clinical adoption. With continued refinement, expanded datasets, and deeper EHR integration, systems of this nature hold the promise of transforming pharmacological decision-making.

FUTURE STUDIES

Future enhancements to the AI Drug Recommendation System could focus on the following directions:

1. Expanded Drug Database:

Extending the classification scope beyond the five drug categories in the UCI dataset to encompass a comprehensive pharmacological formulary.

2. Deep Learning Integration:

Incorporating deep neural network architectures, including transformer-based models, to leverage unstructured clinical text data for enriched drug recommendation context.

3. EHR Integration:

Developing RESTful API connectors for integration with major Electronic Health Record platforms (Epic, Cerner), enabling seamless incorporation into existing clinical workflows.

4. Dosage Quantification:

Extending the system beyond drug class classification to provide specific dosage quantity recommendations based on patient weight, renal function markers, and co-medication profiles.

5. Pharmacogenomics Integration:

Incorporating genetic variant data to enable pharmacogenomics-informed recommendations for medications with established genotype-drug response relationships.

6. Federated Learning:

Implementing federated learning protocols to enable model training across multiple hospital datasets without centralizing sensitive patient data.

7. Ethical AI Framework:

Establishing bias detection and fairness monitoring pipelines to continuously evaluate recommendation equity across demographic subgroups.

8. Multi-Language Support:

Developing multilingual web interfaces to extend accessibility of the drug recommendation system to non-English-speaking clinical environments across Southeast Asia, Latin America, and Sub-Saharan Africa.

9. Real-Time Adverse Drug Reaction Monitoring:

Integrating a post-prescription feedback pipeline that monitors patient outcomes following system-generated drug recommendations, enabling continuous model refinement.

10. Mobile Application Development:

Extending the Flask web application to a cross platform mobile interface using React Native or Flutter frameworks, enabling point-of-care drug recommendation queries.

6.1 Research Limitations and Mitigation Pathways

While the AI Drug Recommendation System demonstrates strong performance on the UCI drug classification dataset, several research limitations must be acknowledged to provide a comprehensive and intellectually honest assessment of the system’s current capabilities. The primary limitation is the restricted size and scope of the UCI drug dataset, which comprises only 200 patient records across five drug classes. While the stratified cross validation protocols employed in this project mitigate overfitting risks, models trained on this scale cannot be assumed to generalize robustly across the full demographic diversity of real-world patient populations — a limitation reinforced quantitatively by the statistical significance analysis in Section 5.2.6, where the 40-record test set proved too small to confirm the dual-ensemble’s accuracy gain with conventional statistical confidence. A second limitation concerns the static nature of the deployed model artifact.

The current implementation serializes the trained model at a fixed point in time, without an automated mechanism for detecting distributional shift in incoming patient data during production operation. A third limitation relates to the exclusion of temporal patient data. The current feature set treats each patient record as an independent, static observation. In reality, drug recommendation in chronic disease management requires longitudinal reasoning that accounts for the trajectory of patient biomarkers over multiple clinical visits.

6.2 Clinical Deployment Roadmap

Transitioning the AI Drug Recommendation System from a research prototype to a clinically deployed decision support tool requires a phased deployment roadmap that systematically addresses regulatory, technical, and organizational readiness dimensions.

1. Extended Dataset Validation:

Collaboration with clinical research partners to acquire multi-institutional patient datasets containing a minimum of 5,000 records across expanded drug categories. Retrospective validation studies would assess model generalization across diverse demographic groups and geographic patient populations, and would also provide a test set large enough to resolve the statistical-power limitation identified in Section 5.2.6.

2. Regulatory Pathway Engagement:

Submission of technical documentation to relevant regulatory bodies under the Software as a Medical Device (SaMD) framework. Documentation would include algorithmic transparency reports, bias assessment studies, and clinical validation evidence.

3. Controlled Pilot Deployment:

Implementation within a controlled clinical pilot environment at one to three partner hospitals, operating as a parallel advisory system alongside conventional prescription workflow.

4. Full-Scale Clinical Integration:

Following successful pilot validation, progressive rollout to broader clinical networks, with EHR API integration, federated learning protocols for continuous model improvement, and pharmacovigilance monitoring pipelines operational. One of the primary future enhancements is the integration of Electronic Health Record (EHR) systems. By connecting the application with hospital databases, the system can automatically retrieve patient information such as medical history, previous prescriptions, laboratory reports, allergies, and chronic diseases. This integration would reduce manual data entry, improve prediction accuracy, and support more personalized treatment recommendations.

The current system utilizes a benchmark dataset with a limited number of patient records. Future versions can be trained using large-scale real-world clinical datasets collected from multiple hospitals and healthcare organizations. A larger and more diverse dataset would improve the model's ability to generalize across different patient populations and medical conditions. Another important enhancement is the incorporation of Deep Learning techniques such as Artificial Neural Networks (ANN), Long Short-Term Memory (LSTM), and Transformer-based models. These advanced algorithms may capture more complex relationships among clinical features and further improve prediction performance, particularly when handling large healthcare datasets.

REFERENCES

- [1] Rajkumar, A., Oren, E., Chen, K., et al. (2018). “Scalable and Accurate Deep Learning with Electronic Health Records.” *npj Digital Medicine*, 1(1), 18.
- [2] Chen, B., Zhang, L., & Wang, Y. (2020). “A Machine Learning Approach to Drug Recommendation Based on Patient Clinical Features.” *Journal of Biomedical Informatics*, 108, 103484.
- [3] Garg, C., Kumar, A., & Pal, S. (2021). “Drug Dosage Prediction Using Ensemble Learning Methods.” *IEEE Access*, 9, 23456–23465.
- [4] Zhang, Q., Liu, W., & Chen, H. (2022). “Explainable AI in Clinical Drug Recommendation: A SHAP-Based Approach.” *Artificial Intelligence in Medicine*, 126, 102258.
- [5] Kumar, R., Singh, P., & Sharma, V. (2021). “Comparative Analysis of Machine Learning Algorithms for Drug Classification in Clinical Decision Support.” *Health Informatics Journal*, 27(3).
- [6] Obermeyer, Z., Powers, B., Vogeli, C., & Mullainathan, S. (2019). “Dissecting racial bias in an algorithm used to manage the health of populations.” *Science*, 366(6464), 447–453.
- [7] Lundberg, S.M., & Lee, S.I. (2017). “A Unified Approach to Interpreting Model Predictions.” *Advances in Neural Information Processing Systems*, 30.
- [8] Breiman, L. (2001). “Random Forests.” *Machine Learning*, 45(1), 5–32.
- [9] Friedman, J.H. (2001). “Greedy Function Approximation: A Gradient Boosting Machine.” *Annals of Statistics*, 29(5), 1189–1232.
- [10] Pedregosa, F., Varoquaux, G., Gramfort, A., et al. (2011). “Scikit-learn: Machine Learning in Python.” *Journal of Machine Learning Research*, 12, 2825–2830.
- [11] UCI Machine Learning Repository. (2020). “Drug Classification Dataset.” University of California, Irvine.
- [12] Topol, E.J. (2019). “High-performance medicine: the convergence of human and artificial intelligence.” *Nature Medicine*, 25(1), 44–56.

- [13] Miotto, R., Wang, F., Wang, S., Jiang, X., & Dudley, J.T. (2018). “Deep learning for healthcare: review, opportunities and challenges.” *Briefings in Bioinformatics*, 19(6), 1236–1246.
- [14] Raza, K. (2019). “Improving the Prediction Accuracy of Heart Disease with Ensemble Learning and Majority Voting Rule.” In *U-Healthcare Monitoring Systems*, 179–196. Academic Press.
- [15] Doshi-Velez, F., & Kim, B. (2017). “Towards a rigorous science of interpretable machine learning.” arXiv preprint arXiv:1702.08608.
- [16] Patel, M., Shah, R., & Gupta, N. (2022). “Flask-Based Clinical Decision Support Web Applications.” *Journal of Medical Internet Research*, 24(5), e34521.
- [17] Richens, J.G., Lee, C.M., & Johri, S. (2020). “Improving the accuracy of medical diagnosis with causal machine learning.” *Nature Communications*, 11, 3923.
- [18] Esteva, A., & Topol, E.J. (2019). “Can AI hurt patients? We need a national agency to find out.” *Nature*, 574, 163–165.
- [19] Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press.
- [20] Rashidi, H.H., Tran, N.K., Betts, E.V., Howell, L.P., & Green, R. (2019). “Artificial Intelligence and Machine Learning in Pathology.” *Academic Pathology*, 6.
- [21] Obermeyer, Z., & Emanuel, E.J. (2016). “Predicting the Future — Big Data, Machine Learning, and Clinical Medicine.” *New England Journal of Medicine*, 375(13), 1216–1219.
- [22] Shortliffe, E.H., & Sepulveda, M.J. (2018). “Clinical Decision Support in the Era of Artificial Intelligence.” *JAMA*, 320(21), 2199–2200.
- [23] Char, D.S., Shah, N.H., & Magnus, D. (2018). “Implementing Machine Learning in Health Care.” *New England Journal of Medicine*, 378(11), 981–983.
- [24] Rudin, C. (2019). “Stop explaining black box machine learning models for high stakes decisions.” *Nature Machine Intelligence*, 1(5), 206–215.
- [25] Jiang, F., Jiang, Y., Zhi, H., et al. (2017). “Artificial intelligence in healthcare: past, present and future.” *Stroke and Vascular Neurology*, 2(4), 230–243.

- [26] Mehrabi, N., Morstatter, F., Saxena, N., Lerman, K., & Galstyan, A. (2021). “A Survey on Bias and Fairness in Machine Learning.” *ACM Computing Surveys*, 54(6), 1–35.
- [27] Rajpurkar, P., Irvin, J., Ball, R.L., et al. (2018). “Deep learning for chest radiograph diagnosis.” *PLOS Medicine*, 15(11).
- [28] Wiens, J., & Shenoy, E.S. (2018). “Machine Learning for Healthcare.” *Clinical Infectious Diseases*, 66(1), 149–153.
- [29] Beam, A.L., & Kohane, I.S. (2018). “Big Data and Machine Learning in Health Care.” *JAMA*, 319(13), 1317–1318.
- [30] Gianfrancesco, M.A., Tamang, S., Yazdany, J., & Schmajuk, G. (2018). “Potential Biases in Machine Learning Algorithms Using Electronic Health Record Data.” *JAMA Internal Medicine*, 178(11), 1544–1547.